

```
In [280]: # Import Libraries
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
import numpy as np
from sklearn.metrics import accuracy_score, confusion_matrix

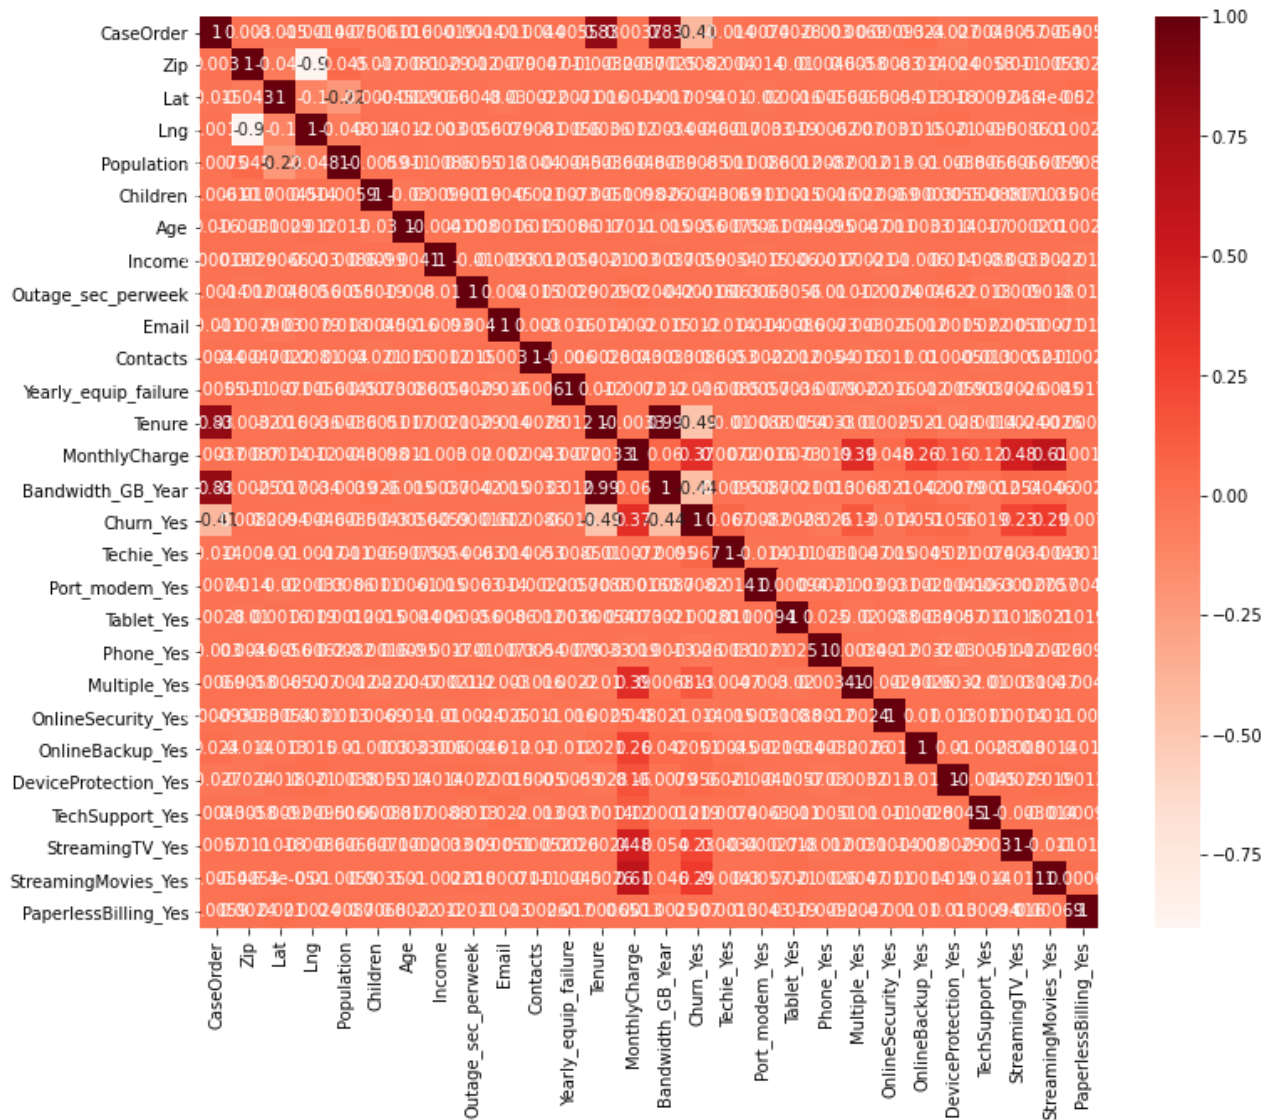
# Import dataset
df = pd.read_csv("C:/Users/hkeim/OneDrive/Documents/School/D208/churn_clean.csv")

In [281]: # Create dummies for binary objects
df=pd.get_dummies(df, columns=['Churn', 'Techie', 'Port_modem', 'Tablet', 'Phone', 'Multiple',
                              'OnlineSecurity', 'OnlineBackup', 'DeviceProte

In [282]: # Drop 'No' dummies
df=df.drop(['Churn_No', 'Techie_No', 'Port_modem_No', 'Tablet_No', 'Phone_No', 'Multiple_No',
            'StreamingTV_No', 'StreamingMovies_No', 'PaperlessBilling_No'], axis=1)

In [283]: #drop unnecessary variables
df=df.drop(['Customer_id', 'Interaction', 'UID', 'City', 'County', 'Job', 'Item1', 'Item2', 'Item5', 'Item6', 'Item7', 'Item8'], axis=1)

In [284]: # Use Pearson Correlation to choose initial model variables
plt.figure(figsize=(12,10))
cor = df.corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.Reds)
plt.show()
```



```
In [285]: # Create dataframe with initial variables
df=df.filter(items=['CaseOrder', 'Tenure', 'MonthlyCharge', 'Bandwidth_GB_Year', 'Churn_Yes',
```

```
In [286]: # Summary statistics
print(df.describe())
```

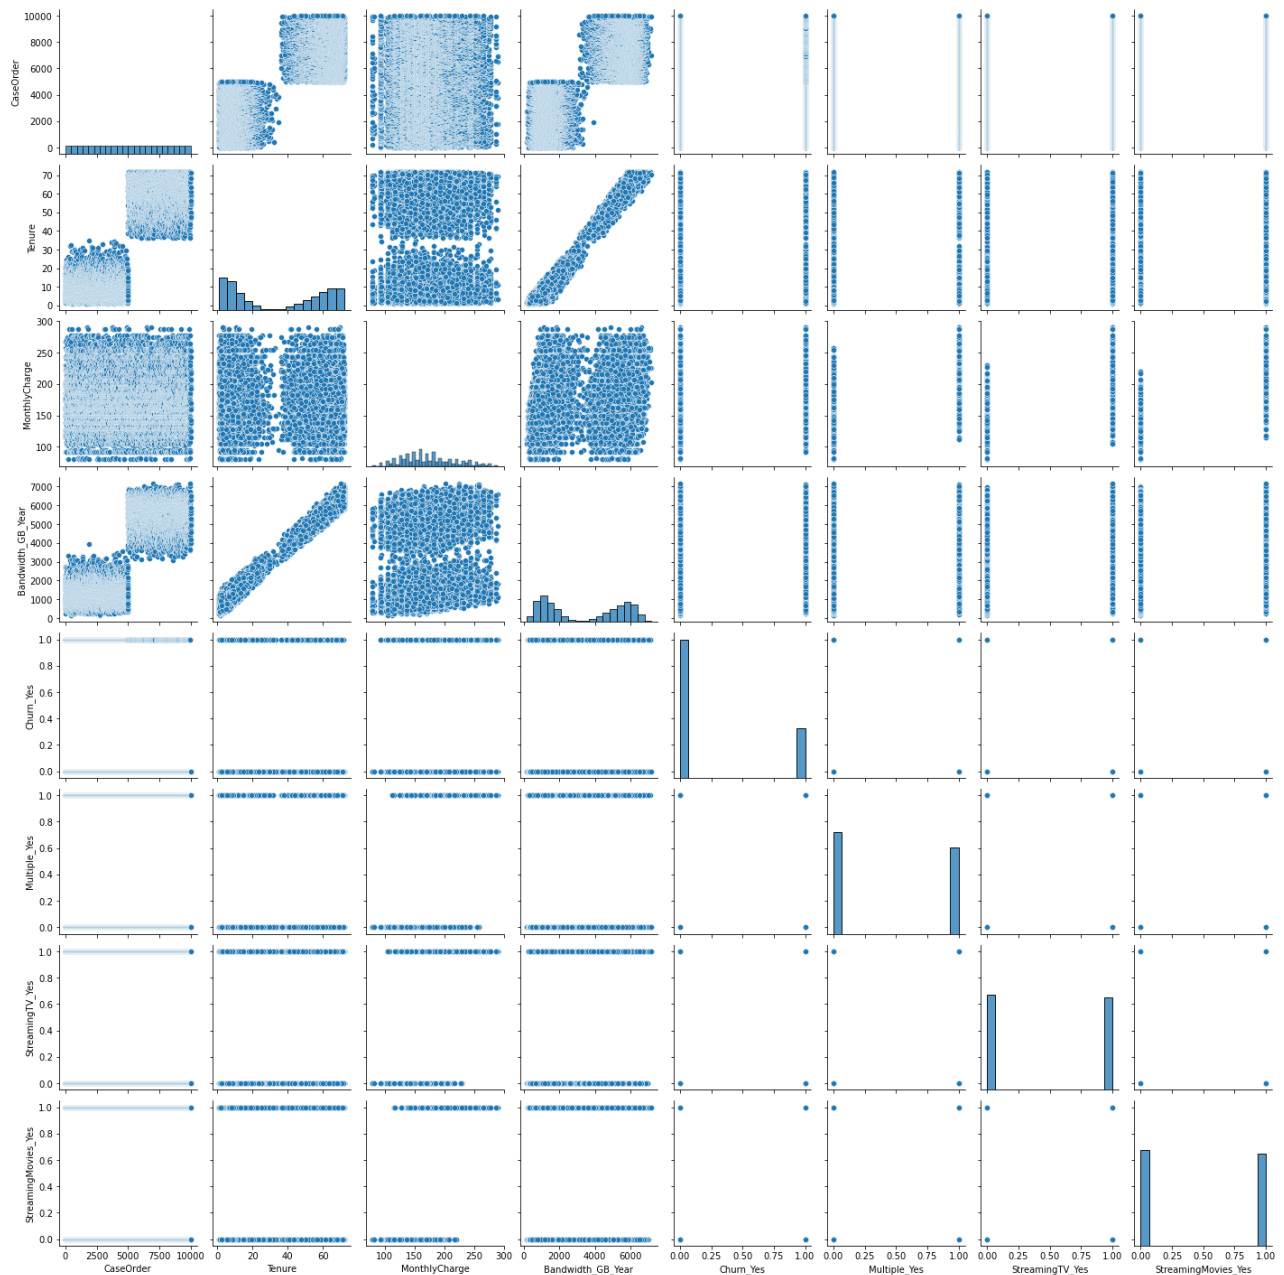
	CaseOrder	Tenure	MonthlyCharge	Bandwidth_GB_Year	\
count	10000.00000	10000.00000	10000.00000	10000.00000	
mean	5000.50000	34.526188	172.624816	3392.341550	
std	2886.89568	26.443063	42.943094	2185.294852	
min	1.00000	1.000259	79.978860	155.506715	
25%	2500.75000	7.917694	139.979239	1236.470827	
50%	5000.50000	35.430507	167.484700	3279.536903	
75%	7500.25000	61.479795	200.734725	5586.141370	
max	10000.00000	71.999280	290.160419	7158.981530	

	Churn_Yes	Multiple_Yes	StreamingTV_Yes	StreamingMovies_Yes
count	10000.00000	10000.00000	10000.00000	10000.00000
mean	0.265000	0.460800	0.492900	0.489000
std	0.441355	0.498486	0.499975	0.499904
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000

75%	1.000000	1.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

```
In [287]: # Univariate and bivariate visualizations
sns.pairplot(df)
```

```
Out[287]: <seaborn.axisgrid.PairGrid at 0x2a7451fc6d0>
```



```
In [288]: # Defining the independent variables
X=df[['CaseOrder', 'Tenure', 'MonthlyCharge', 'Bandwidth_GB_Year', 'Multiple_Yes', 'Streaming
```

```
In [289]: # Get initial model intercept
X = sm.add_constant(X)
```

```
In [290]: # Defining the dependent Variable
y=df['Churn_Yes']
```

In [291]:

```
# Initial logistic regression model
model = sm.Logit(endog=y, exog=X).fit()
print(model.summary())
```

Optimization terminated successfully.
Current function value: inf
Iterations 8

```
Logit Regression Results
=====
Dep. Variable:          Churn_Yes    No. Observations:          10000
Model:                  Logit        Df Residuals:              9992
Method:                  MLE         Df Model:                  7
Date:                   Sat, 09 Oct 2021    Pseudo R-squ.:            inf
Time:                   21:46:55          Log-Likelihood:           -inf
Converged:              True           LL-Null:                  0.0000
Covariance Type:        nonrobust        LLR p-value:              1.000
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-3.8150	0.193	-19.784	0.000	-4.193	-3.437
CaseOrder	6.086e-06	2.01e-05	0.302	0.762	-3.34e-05	4.55e-05
Tenure	-0.2347	0.012	-19.684	0.000	-0.258	-0.211
MonthlyCharge	0.0104	0.002	6.730	0.000	0.007	0.013
Bandwidth_GB_Year	0.0019	0.000	13.694	0.000	0.002	0.002
Multiple_Yes	0.6336	0.082	7.766	0.000	0.474	0.793
StreamingTV_Yes	1.1356	0.099	11.498	0.000	0.942	1.329
StreamingMovies_Yes	1.4812	0.109	13.642	0.000	1.268	1.694

```
=====
```

In [292]:

```
# Initial model odds ratios
np.exp(model.params)
```

```
Out[292]: const          0.022037
CaseOrder      1.000006
Tenure         0.790815
MonthlyCharge  1.010412
Bandwidth_GB_Year 1.001879
Multiple_Yes   1.884349
StreamingTV_Yes 3.113014
StreamingMovies_Yes 4.398402
dtype: float64
```

In [293]:

```
# Initial model predictions
pred = model.predict(exog=X)
pred.head()
```

```
Out[293]: 0    0.390642
1    0.960092
2    0.528901
3    0.200313
4    0.267200
dtype: float64
```

In [294]:

```
# Initial model rounded predictions
round(pred)
```

```
Out[294]: 0    0.0
1    1.0
2    1.0
3    0.0
4    0.0
...
```

```
9995    0.0
9996    0.0
9997    0.0
9998    0.0
9999    0.0
Length: 10000, dtype: float64
```

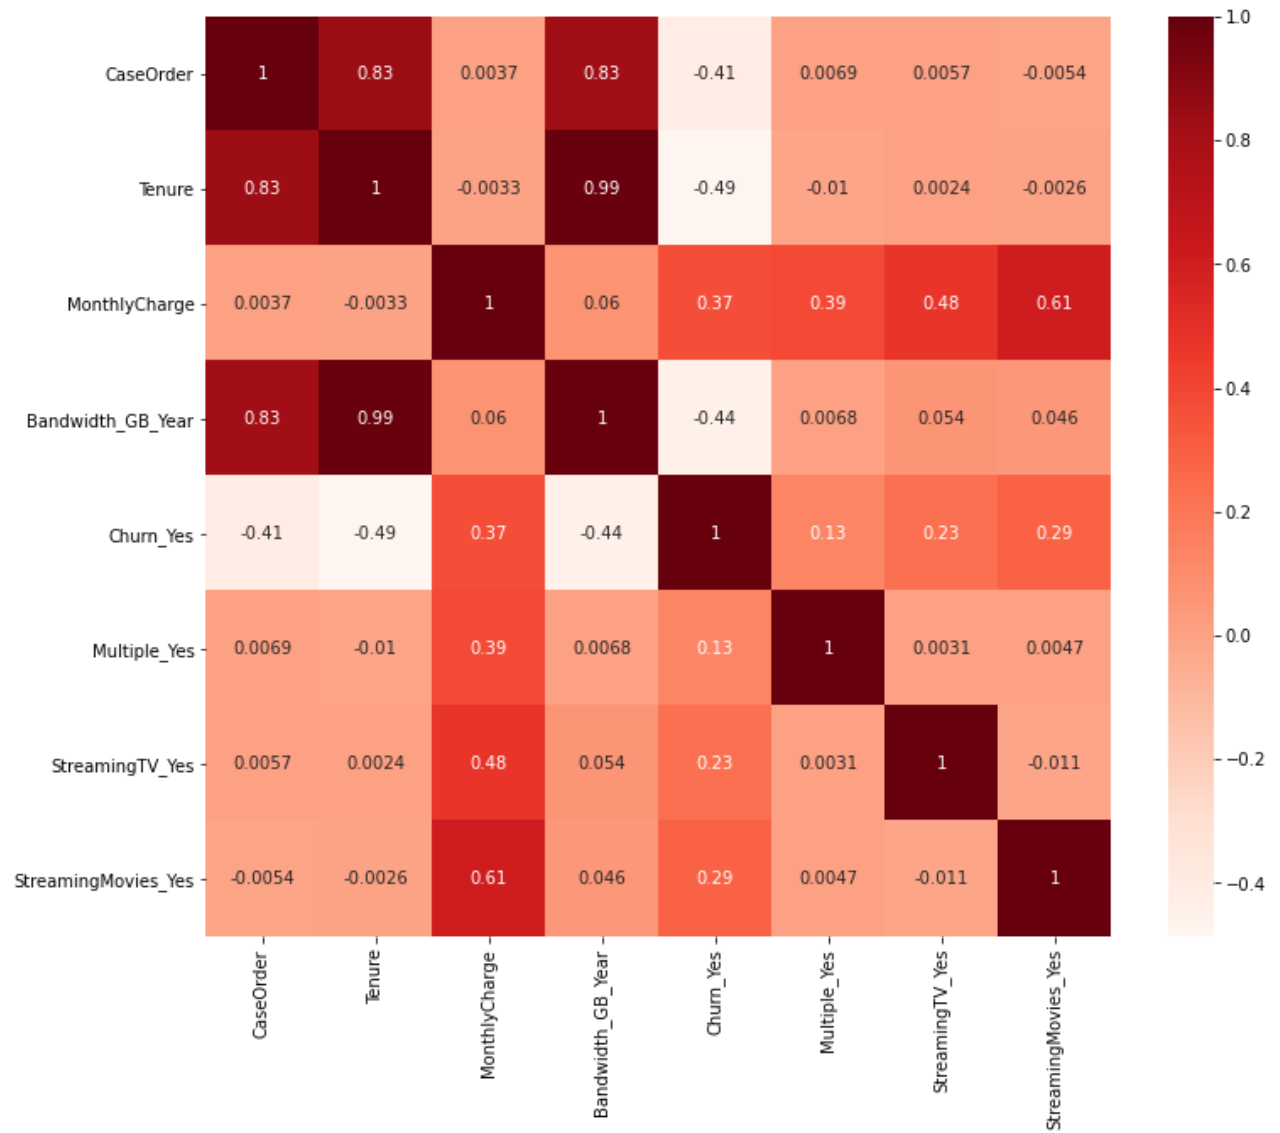
```
In [295]: # Initial model confusion matrix
          confusion_matrix(y_true=list(y), y_pred=list(round(pred)))
```

```
Out[295]: array([[6799,  551],
                 [ 944, 1706]], dtype=int64)
```

```
In [296]: # Initial model accuracy of the fitted model
          accuracy_score(y_true=list(y), y_pred=list(round(pred)))
```

```
Out[296]: 0.8505
```

```
In [297]: # Use Pearson Correlation to choose reduced model variables
          plt.figure(figsize=(12,10))
          cor = df.corr()
          sns.heatmap(cor, annot=True, cmap=plt.cm.Reds)
          plt.show()
```



```
In [298]: # Pairwise correlation analysis
X.corr()
```

Out[298]:

	const	CaseOrder	Tenure	MonthlyCharge	Bandwidth_GB_Year	Multiple_Yes	Streamin
const	NaN	NaN	NaN	NaN	NaN	NaN	
CaseOrder	NaN	1.000000	0.832550	0.003677	0.825561	0.006915	
Tenure	NaN	0.832550	1.000000	-0.003337	0.991495	-0.010422	
MonthlyCharge	NaN	0.003677	-0.003337	1.000000	0.060406	0.385979	
Bandwidth_GB_Year	NaN	0.825561	0.991495	0.060406	1.000000	0.006823	
Multiple_Yes	NaN	0.006915	-0.010422	0.385979	0.006823	1.000000	
StreamingTV_Yes	NaN	0.005690	0.002440	0.482312	0.054314	0.003097	
StreamingMovies_Yes	NaN	-0.005353	-0.002574	0.608115	0.045600	0.004691	-

```
In [299]: # Defining the reduced independent variables
Xr=df[['MonthlyCharge', 'Bandwidth_GB_Year', 'Multiple_Yes', 'StreamingTV_Yes']]
```

```
In [300]: # Get reduced model intercept
Xr = sm.add_constant(Xr)
```

```
In [301]: # Reduced Logistic regression model
modelr = sm.Logit(endog=y, exog=Xr).fit()
print(modelr.summary())
```

Optimization terminated successfully.

Current function value: inf

Iterations 7

Logit Regression Results

```
=====
Dep. Variable:          Churn_Yes    No. Observations:          10000
Model:                  Logit        Df Residuals:              9995
Method:                  MLE         Df Model:                  4
Date:                   Sat, 09 Oct 2021    Pseudo R-squ.:            inf
Time:                   21:48:25    Log-Likelihood:           -inf
converged:              True         LL-Null:                  0.0000
Covariance Type:        nonrobust    LLR p-value:              1.000
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-4.7256	0.149	-31.783	0.000	-5.017	-4.434
MonthlyCharge	0.0311	0.001	30.962	0.000	0.029	0.033
Bandwidth_GB_Year	-0.0008	2e-05	-41.530	0.000	-0.001	-0.001
Multiple_Yes	0.0184	0.065	0.284	0.777	-0.109	0.145
StreamingTV_Yes	0.5945	0.069	8.635	0.000	0.460	0.729

```
=====
```

```
In [302]: # Reduced model odds ratios
np.exp(modelr.params)
```

```
Out[302]: const          0.008866
MonthlyCharge    1.031545
Bandwidth_GB_Year 0.999171
Multiple_Yes     1.018540
StreamingTV_Yes  1.812107
dtype: float64
```

```
In [303]: # Reduced model predictions
pred = modelr.predict(exog=Xr)
pred.head()
```

```
Out[303]: 0    0.470017
1    0.940376
2    0.190846
3    0.099609
4    0.574606
dtype: float64
```

```
In [304]: # Reduced model rounded predictions
round(predr)
```

```
Out[304]: 0    0.0
1    1.0
2    0.0
3    0.0
4    1.0
...
9995  0.0
```

```
9996    0.0
9997    0.0
9998    0.0
9999    0.0
Length: 10000, dtype: float64
```

```
In [305]: # Reduced model confusion matrix
          confusion_matrix(y_true=list(y), y_pred=list(round(predr)))
```

```
Out[305]: array([[6744,  606],
                 [1080, 1570]], dtype=int64)
```

```
In [306]: # Reduced model accuracy of the fitted model
          accuracy_score(y_true=list(y), y_pred=list(round(predr)))
```

```
Out[306]: 0.8314
```

```
In [307]: # Export regression data set
          df.to_csv("C:/Users/hkeim/OneDrive/Documents/School/D208/Keim D208 Task Two Prepared Data.csv")
```

```
In [ ]:
```